

```
In [1]: cd /content/drive/My Drive/Xray report generation

/content/drive/My Drive/Xray report generation
```

```
In [2]: ls

AttentionModel.ipynb          Medical.csv
Basic_Model.ipynb            model_1_plot.png
brucechou1983_CheXNet_Keras_0.3.0_weights.h5  NLMCXR_png/
checkpoints/                  'Old Files'/
Chexnetmodel/                 o_test.pkl
'Copy of AttentionModel.ipynb' preprocessed.csv
'Copy of SimpleModel.ipynb'   saved_model/
dup_preprocessed.csv          simple_encoder_decoder_plot.png
ecgen-radiology/              SimpleModel.ipynb
image_feature_vector1.pkl     tokenizer1.pkl
indiana_projections.csv       Untitled0.ipynb
i_test.pkl                    Untitled1.ipynb
logs/                          Untitled2.ipynb
logs7/
```

```
In [3]: import re
import pandas as pd
import matplotlib.pyplot as plt
import string
from sklearn.utils import shuffle
from sklearn.model_selection import train_test_split
from tqdm import tqdm
import numpy as np
import tensorflow as tf
import nltk
```

```
In [ ]: df = pd.read_csv('Medical.csv')
```

```
In [ ]: df.head()
```

Out[74]:

	image_name	image_caption	comparison	indication	findings	impression
0	CXR1_1_IM-0001-3001.png,CXR1_1_IM-0001-4001.png	xray chest pa and lateral	none	positive tb test	the cardiac silhouette and mediastinum size ar...	normal chest x
1	CXR10_IM-0002-2001.png,CXR10_IM-0002-1001.png	pa and lateral chest x	chest radiographs	male chest pain	the cardiomeastinal silhouette is within nor...	no acute cardiopulmonary process
2	CXR100_IM-0002-1001.png,CXR100_IM-0002-2001.png	chest v frontallateral pm	none	no indication	both lungs are clear and expanded heart and me...	no active disease
3	CXR1000_IM-0003-1001.png,CXR1000_IM-0003-2001....	pa and lateral chest x	pa and lateral chest radiographs	male	there is increased opacity within the right up...	increased opacity in the right upper lobe with...
4	CXR1001_IM-0004-1001.png,CXR1001_IM-0004-1002.png	chest v frontallateral pm	none	dyspnea subjective fevers arthritis immigrant ...	interstitial markings are diffusely prominent ...	diffuse fibrosis no visible focal acute disease

```
In [ ]: for item in tqdm(df.iterrows()):
        print(item[1])
        break
```

0it [00:00, ?it/s]

```
image_name          CXR1_1_IM-0001-3001.png,CXR1_1_IM-0001-4001.png
image_caption              xray chest pa and lateral
comparison              none
indication              positive tb test
findings              the cardiac silhouette and mediastinum size ar...
impression              normal chest x
findings_count              33
impression_count              3
image_count              2
Name: 0, dtype: object
```

```
In [ ]: df_projections = pd.read_csv('indiana_projections.csv')
```

```
In [ ]: df_projections.head()
```

```
Out[32]:
```

	uid	filename	projection
0	1	1_IM-0001-4001.dcm.png	Frontal
1	1	1_IM-0001-3001.dcm.png	Lateral
2	2	2_IM-0652-1001.dcm.png	Frontal
3	2	2_IM-0652-2001.dcm.png	Lateral
4	3	3_IM-1384-1001.dcm.png	Frontal

```
In [ ]: img_path = 'NLM CXR_png/'
```

```
In [ ]: # to find the matching pattern in a string - https://thepythonguru.com/python-regular-expressions/
def find_fr_lr_images(li):
    """
    This funct is used to find the lateral images and frontal images from our list
    """
    img_list = []
    last_imag = ""
    for img in li:
        projections = df_projections[df_projections['filename'].str.contains(re.search('Lateral', img))]
        if 'Lateral' == projections:
            last_imag = img
        else:
            img_list.append(img)
    return img_list, last_imag
```

```

In [ ]: #Creating structured data from raw xml files
columns = ["image_1", "image_2", "impression"]
df_1 = pd.DataFrame(columns = columns)
columns = ["image_1", "image_2", "impression"]
df_dup = pd.DataFrame(columns = columns)
no_lateral = 0
for item in tqdm(df.iterrows()):
    l = item[1]['image_name'].split(',')
    if len(l) > 2:
        li, last_img = find_fr_lr_images(l)
        if last_img == "":
            no_lateral +=1
            li, last_img = li[:-1], li[-1]
        for i in li:
            image_1 = i
            image_2 = last_img
            df_1 = df_1.append(pd.Series([image_1, image_2, item[1]['findings']],
    elif len(l) == 2:
        image_1 = l[0]
        image_2 = l[1]
        df_1 = df_1.append(pd.Series([image_1, image_2, item[1]['findings']], inc
    elif len(l) == 1:
        #creating duplicate dataframe separately to keep it in all dataset train
        df_dup = df_dup.append(pd.Series([l[0], l[0], item[1]['findings']], index
print("Total Report without Lateral images {}".format(no_lateral))

```

1130it [00:04, 264.40it/s]/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:10: DeprecationWarning: The truth value of an empty array is ambiguous. Returning False, but in future this will result in an error. Use `array.size > 0` to check that an array is not empty.

# Remove the CWD from sys.path while we load stuff.

3851it [00:15, 255.66it/s]

Total Report without Lateral images 1

```

In [ ]: df_1.shape

```

Out[58]: (3532, 3)

```
In [ ]: df_1.head()
```

```
Out[59]:
```

	image_1	image_2	impression
0	CXR1_1_IM-0001-3001.png	CXR1_1_IM-0001-4001.png	cardiac silhouette mediastinum size within nor...
1	CXR10_IM-0002-1001.png	CXR10_IM-0002-2001.png	cardiomediastinal silhouette within normal lim...
2	CXR100_IM-0002-1001.png	CXR100_IM-0002-2001.png	lungs clear expanded heart mediastinum normal
3	CXR1000_IM-0003-1001.png	CXR1000_IM-0003-2001.png	increased opacity within right upper lobe poss...
4	CXR1000_IM-0003-3001.png	CXR1000_IM-0003-2001.png	increased opacity within right upper lobe poss...

```
In [ ]: df_dup.shape
```

```
Out[60]: (446, 3)
```

```
In [ ]: df_dup.head()
```

```
Out[61]:
```

	image_1	image_2	impression
0	CXR1003_IM-0005-2002.png	CXR1003_IM-0005-2002.png	heart size pulmonary vascularity appear within...
1	CXR1012_IM-0013-1001.png	CXR1012_IM-0013-1001.png	lung volumes low bibasilar consolidation bilat...
2	CXR1024_IM-0019-1001.png	CXR1024_IM-0019-1001.png	heart size normal lungs clear no focal air spa...
3	CXR1026_IM-0021-2002.png	CXR1026_IM-0021-2002.png	lungs appear clear heart pulmonary appear norm...
4	CXR1029_IM-0022-1001.png	CXR1029_IM-0022-1001.png	NaN

Adding Start and end tokens

```
In [ ]: def add_start_end_token(data):
    # Combining all the above students
    preprocessed_reviews_eng = []

    # tqdm is for printing the status bar
    for sentence in tqdm(data.values):
        sentence = '<start> ' + sentence + ' <end>'
        preprocessed_reviews_eng.append(sentence.strip())
    return preprocessed_reviews_eng
```

```
In [ ]: df_1['impression'] = add_start_end_token(df_1['impression'].astype(str))
```

```
100%|██████████| 3532/3532 [00:00<00:00, 514277.64it/s]
```

```
In [ ]: df_dup['impression'] = add_start_end_token(df_dup['impression'].astype(str))
```

```
100%|██████████| 446/446 [00:00<00:00, 376238.85it/s]
```

```
In [ ]: df_1.head()
```

```
Out[65]:
```

	image_1	image_2	impression
0	CXR1_1_IM-0001-3001.png	CXR1_1_IM-0001-4001.png	<start> cardiac silhouette mediastinum size wi...
1	CXR10_IM-0002-1001.png	CXR10_IM-0002-2001.png	<start> cardiomedastinal silhouette within no...
2	CXR100_IM-0002-1001.png	CXR100_IM-0002-2001.png	<start> lungs clear expanded heart mediastinum...
3	CXR1000_IM-0003-1001.png	CXR1000_IM-0003-2001.png	<start> increased opacity within right upper l...
4	CXR1000_IM-0003-3001.png	CXR1000_IM-0003-2001.png	<start> increased opacity within right upper l...

```
In [ ]: df_dup.head()
```

```
Out[66]:
```

	image_1	image_2	impression
0	CXR1003_IM-0005-2002.png	CXR1003_IM-0005-2002.png	<start> heart size pulmonary vascularity appea...
1	CXR1012_IM-0013-1001.png	CXR1012_IM-0013-1001.png	<start> lung volumes low bibasilar consolidati...
2	CXR1024_IM-0019-1001.png	CXR1024_IM-0019-1001.png	<start> heart size normal lungs clear no focal...
3	CXR1026_IM-0021-2002.png	CXR1026_IM-0021-2002.png	<start> lungs appear clear heart pulmonary app...
4	CXR1029_IM-0022-1001.png	CXR1029_IM-0022-1001.png	<start> nan <end>

```
In [ ]: df_1.to_csv('preprocessed1.csv')
```

```
In [ ]: df_dup.to_csv('dup_preprocessed1.csv')
```

```
In [ ]: df_1 = pd.read_csv('preprocessed.csv')
```

```
In [ ]: df_1.head()
```

```
Out[76]:
```

	Unnamed: 0	image_1	image_2	impression
0	0	CXR1_1_IM-0001-3001.png	CXR1_1_IM-0001-4001.png	<start> normal chest x <end>
1	1	CXR10_IM-0002-2001.png	CXR10_IM-0002-1001.png	<start> no acute cardiopulmonary process <end>
2	2	CXR100_IM-0002-1001.png	CXR100_IM-0002-2001.png	<start> no active disease <end>
3	3	CXR1000_IM-0003-1001.png	CXR1000_IM-0003-2001.png	<start> increased opacity in the right upper l...
4	4	CXR1000_IM-0003-3001.png	CXR1000_IM-0003-2001.png	<start> increased opacity in the right upper l...

```
In [ ]: df_dup = pd.read_csv('dup_preprocessed.csv')
```

```
In [ ]: df_dup.head()
```

```
Out[78]:
```

	Unnamed: 0	image_1	image_2	impression
0	0	CXR1003_IM-0005-2002.png	CXR1003_IM-0005-2002.png	<start> retrocardiac soft tissue density the a...
1	1	CXR1012_IM-0013-1001.png	CXR1012_IM-0013-1001.png	<start> bibasilar airspace disease and bilater...
2	2	CXR1024_IM-0019-1001.png	CXR1024_IM-0019-1001.png	<start> no acute abnormality <end>
3	3	CXR1026_IM-0021-2002.png	CXR1026_IM-0021-2002.png	<start> no acute cardiopulmonary disease <end>
4	4	CXR1029_IM-0022-1001.png	CXR1029_IM-0022-1001.png	<start> no pneumonia heart size normal scolios...

```
In [ ]: df_1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3532 entries, 0 to 3531
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0   3532 non-null   int64
1   image_1      3532 non-null   object
2   image_2      3532 non-null   object
3   impression   3532 non-null   object
dtypes: int64(1), object(3)
memory usage: 110.5+ KB
```

```
In [ ]: # df_1.drop("Unnamed: 0",axis='columns',inplace=True)
```

```
In [ ]: image_names = []
for img in tqdm(df['image_name'].str.split(',')):

    for i in range(len(img)):
        image_names.append(img[i])
```

100%|██████████| 3851/3851 [00:00<00:00, 560278.36it/s]

```
In [ ]: image_names[0:4]
```

```
Out[82]: ['CXR1_1_IM-0001-3001.png',
          'CXR1_1_IM-0001-4001.png',
          'CXR10_IM-0002-2001.png',
          'CXR10_IM-0002-1001.png']
```

```
In [ ]: import tensorflow as tf
from tensorflow.keras.applications import densenet
from tensorflow.keras.applications.densenet import preprocess_input
from tensorflow.keras.layers import Dense, Dropout, Input, Conv2D
from tensorflow.keras.models import Model
```

```
In [ ]: ## the input layer shape is (None,None,None,3)
## the output layer is (None,2048)
```

```
In [ ]: chex = densenet.DenseNet121(include_top=False,weights=None,input_shape=(224,224,3))
X = chex.output
X = Dense(14,activation='sigmoid',name='predictions')(X)
model = Model(inputs=chex.input,outputs=X)
model.load_weights('brucechou1983_CheXNet_Keras_0.3.0_weights.h5')
image_features_model = Model(inputs=model.input,outputs= model.layers[-2].output)
```

```
In [ ]: import pickle
with open('image_feature_vector1.pkl','rb') as f:
    image_feature_vector_1 = pickle.load(f)
```

```
In [ ]: print(image_feature_vector_1[1])
```

```
tf.Tensor(
[[5.0062913e-04 2.0340595e-03 1.2797897e-03 ... 9.1456938e-01
 9.2038012e-01 7.4285048e-01]], shape=(1, 1024), dtype=float32)
```

Train, Validation and Test split

```
In [ ]: i_train,input_test,o_train,output_test = train_test_split(df_1[['image_1','image_2'],
input_train,input_validate,output_train,output_validate = train_test_split(i_train
```



```
In [ ]: input_train.shape,input_validate.shape,input_test.shape,output_train.shape,output_test.shape
```

```
Out[88]: ((2542, 2), (636, 2), (354, 2), (2542,), (636,), (354,))
```

```
In [ ]: i_train_dup,input_test_dup,o_train_dup,output_test_dup = train_test_split(df_dup,
input_train_dup,input_validate_dup,output_train_dup,output_validate_dup = train_test_split
```

```
In [ ]: input_train_dup.shape,input_validate_dup.shape,input_test_dup.shape,output_train_dup.shape,output_test_dup.shape
```

```
Out[90]: ((320, 2), (81, 2), (45, 2), (320,), (81,), (45,))
```

Appending the duplicate data and the original data

```
In [ ]: i_train = np.append(input_train,input_train_dup,axis=0)
o_train = np.append(output_train,output_train_dup,axis=0)
i_validate = np.append(input_validate,input_validate_dup,axis=0)
o_validate = np.append(output_validate,output_validate_dup,axis=0)
i_test = np.append(input_test,input_test_dup,axis=0)
o_test = np.append(output_test,output_test_dup,axis=0)
```

```
In [ ]: i_train[0]
```

```
Out[92]: array(['CXR914_IM-2417-1001.png', 'CXR914_IM-2417-3001.png'], dtype=object)
```

```
In [ ]: #https://datascience.stackexchange.com/questions/24511/why-should-the-data-be-shuffled
# Here we will be shuffling the data
for i in range(5):
    i_train,o_train = shuffle(i_train,o_train,random_state=15)
    i_validate,o_validate = shuffle(i_validate,o_validate,random_state=15)
    i_test,o_test = shuffle(i_test,o_test,random_state=15)
```

## TEXT TOKENIZATION

```
In [ ]: from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

max_length_output = 60

tokenizer = Tokenizer(oov_token="<unk>", filters='!"#$%&()*+.,-/:;=?@[\\]^_`{|}~')
tokenizer.fit_on_texts(o_train)
text_train = tokenizer.texts_to_sequences(o_train)
text_validation = tokenizer.texts_to_sequences(o_validate)
text_test = tokenizer.texts_to_sequences(o_test)
```

```
In [ ]: import pickle

with open('tokenizer1.pkl','wb') as f:
    pickle.dump(tokenizer,f)
```

```
In [ ]: dictionary = tokenizer.word_index
```

```
In [ ]: word2idx = {}
idx2word = {}
for k,v in dictionary.items():
    word2idx[k] = v # put the word value for the particular key(index)
    idx2word[v] = k # put the particular index for the word
```

```
In [ ]: vocabulary_size = len(word2idx)+1
print(vocabulary_size)
```

1339

```
In [ ]: print("==== Top 6 Word and its Index =====")
list(dictionary.items())[:6]
```

==== Top 6 Word and its Index =====

```
Out[99]: [('<unk>', 1),
          ('<start>', 2),
          ('<end>', 3),
          ('no', 4),
          ('acute', 5),
          ('cardiopulmonary', 6)]
```

```
In [ ]: ## Padding the text
text_train_output = pad_sequences(text_train,maxlen=60,dtype='int32',padding='post')
text_validation_output = pad_sequences(text_validation,maxlen=60,dtype='int32',padding='post')
text_test_output = pad_sequences(text_test,maxlen=60,dtype='int32',padding='post')
```

```
In [ ]: text_train_output.shape
```

```
Out[101]: (2862, 60)
```

Creating Tensorflow Dataset

```
In [ ]: # converting the numpy image files into tensorformats

def multi_image(img,imp):
    """
    This function will be taking in the imgaes we have and converting into tensor format
    """
    return tf.convert_to_tensor([image_feature_vector_1[image_names.index(img[0]).decode('utf-8')]])
```

```
In [ ]: dataset_train = tf.data.Dataset.from_tensor_slices((i_train,text_train_output))
```

```
In [ ]: dataset_validation = tf.data.Dataset.from_tensor_slices((i_validate,text_validate))
```

```
In [ ]: dataset_train = dataset_train.map(lambda item1,item2: tf.numpy_function(multi_image_processing, [item1, item2], tf.float32))
dataset_validation = dataset_validation.map(lambda item1,item2: tf.numpy_function(multi_image_processing, [item1, item2], tf.float32))
```

```
In [ ]: BATCH_SIZE = 32
BUFFER_SIZE = 1000
embedding_dimension = 256
units = 128 # output
```

```
In [ ]: # Shuffle and Batch
dataset_train = dataset_train.shuffle(BUFFER_SIZE).batch(BATCH_SIZE)
dataset_train = dataset_train.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)

#Shuffle and Batch
dataset_validation = dataset_validation.shuffle(BUFFER_SIZE).batch(BATCH_SIZE)
dataset_validation = dataset_validation.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
```

```
In [ ]: # input_shape is (2,1,2048)
# vocab_size = 1337
```

Encoder Decoder Model

```
In [ ]: class BahdanauAttention(tf.keras.Model):
    def __init__(self, units):
        super(BahdanauAttention, self).__init__()
        self.w1 = tf.keras.layers.Dense(units)
        self.w2 = tf.keras.layers.Dense(units)
        self.v = tf.keras.layers.Dense(1)

    def call(self, features, hidden):

        ## features(Encoder output) shape == (batch_size, 64, embedding_dim)
        ## hidden shape == (batch_size, hidden_size)
        ## hidden_with_time_axis == (batch_size, 1, hidden_size)

        hidden_with_time_axis = tf.expand_dims(hidden, 1)

        # attention_hidden_layer shape == (batch_size, 49, units)

        attention_hidden_layer = (tf.nn.tanh(self.w1(features) + self.w2(hidden_with_time_axis)))

        # score shape == (batch_size, 64, 1)
        # This gives an unnormalized score for each image feature.
        score = self.v(attention_hidden_layer)
        attention_weights = tf.nn.softmax(score, axis=1)

        # context_vector shape after sum == (batch_size, hidden_size)
        context_vector = attention_weights * features
        context_vector = tf.reduce_sum(context_vector, axis=1)

        return context_vector, attention_weights
```

```
In [ ]: class Encoder(tf.keras.Model):
    def __init__(self, embedding_dim):
        super(Encoder, self).__init__()
        self.dense = tf.keras.layers.Dense(embedding_dim, activation='relu', kernel_initializer='he_normal')
        self.concat = tf.keras.layers.Concatenate()
    def call(self, x):

        encoder_concat = self.concat([x[:, 0], x[:, 1]])
        x = self.dense(encoder_concat)
        x = tf.nn.relu(x)
        return x
```

```

In [ ]: class RNN_Decoder(tf.keras.Model):
    def __init__(self, embedding_dim, units, vocab_size):
        super(RNN_Decoder, self).__init__()
        self.units = units

        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
        self.gru = tf.keras.layers.GRU(self.units, return_sequences=True, return_state=True)
        self.fc1 = tf.keras.layers.Dense(self.units)
        self.fc2 = tf.keras.layers.Dense(vocab_size)
        self.attention = BahdanauAttention(self.units)

    def call(self, x, features, hidden):

        #defining attention as a separate model
        context_vector, attention_weights = self.attention(features, hidden)
        # shape of x after embedding (batch_size, 1, embedding_dimension)
        x = self.embedding(x)
        # shape of x after concatenation (batch_size, 1, embedding_dimension + hidden_size)
        x = tf.concat([tf.expand_dims(context_vector, 1), x], axis=-1)

        # passing the concatenated vector to the GRU

        output, state = self.gru(x)

        # shape == (batch_size, max_length, hidden_size)
        x = self.fc1(output)

        # shape == (batch_size * max_length, hidden_size)
        x = tf.reshape(x, (-1, x.shape[2]))

        # output shape == (batch_size * max_length, vocab)

        x = self.fc2(x)

        return x, state, attention_weights

```

In [ ]:

```

In [ ]: encoder = Encoder(embedding_dimension)
        decoder = RNN_Decoder(embedding_dimension, units, vocabulary_size)

```

```
In [ ]: # modelling loss and accuracy for gradient tape based on this https://www.tensorflow.org/api\_guides/python/keras

optimizer = tf.keras.optimizers.Adam()
loss_obj = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True, reduction='none')
acc_obj = tf.keras.metrics.SparseCategoricalAccuracy()

def loss_func(real, pred):
    mask = tf.math.logical_not(tf.math.equal(real, 0))
    loss_ = loss_obj(real, pred)

    mask = tf.cast(mask, dtype=loss_.dtype)
    loss_ *= mask

    return tf.reduce_mean(loss_)

def acc_func(real, pred):
    acc_f = acc_obj(real, pred)
    return tf.reduce_mean(acc_f)
```

```
In [ ]: checkpoint_path = "./checkpoints/train"
ckpt = tf.train.Checkpoint(encoder=encoder, decoder=decoder, optimizer=optimizer)

ckpt_manager = tf.train.CheckpointManager(ckpt, checkpoint_path, max_to_keep=5)
```

```
In [ ]: start_epoch = 0
if ckpt_manager.latest_checkpoint:
    start_epoch = int(ckpt_manager.latest_checkpoint.split('-')[-1])
    # restoring the latest checkpoint in checkpoint_path
    ckpt.restore(ckpt_manager.latest_checkpoint)
```

```
In [ ]: import datetime
current_time = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
train_log_dir = 'logs7/Basic_gradient_tape/' + current_time + '/train'
validation_log_dir = 'logs7/Basic_gradient_tape/' + current_time + '/test'
train_summary_writer = tf.summary.create_file_writer(train_log_dir)
validation_summary_writer = tf.summary.create_file_writer(validation_log_dir)
```

MODEL Training

```

In [ ]: @tf.function
def train_step(img_tensor, target):
    """Train Step"""
    loss = 0
    accuracy = 0
    hidden = tf.zeros((target.shape[0], units))
    dec_input = tf.expand_dims([tokenizer.word_index['<start>']] * target.shape[0], 1)

    with tf.GradientTape() as tape:
        features = encoder(img_tensor)

        for i in range(1, target.shape[1]):
            # passing the features through the decoder
            predictions, hidden, _ = decoder(dec_input, features, hidden)
            loss += loss_func(target[:, i], predictions)
            accuracy += acc_func(target[:, i], predictions)

            # using teacher forcing
            dec_input = tf.expand_dims(target[:, i], 1)

    total_loss = (loss / int(target.shape[1]))
    total_accuracy = (accuracy / int(target.shape[1]))

    trainable_variables = encoder.trainable_variables + decoder.trainable_variables
    gradients = tape.gradient(loss, trainable_variables)

    optimizer.apply_gradients(zip(gradients, trainable_variables))

    return loss, total_loss, total_accuracy

@tf.function
def val_step(img_tensor, target):
    """Train Step"""
    loss_val = 0
    accuracy_val = 0
    hidden = tf.zeros((target.shape[0], units))
    dec_input = tf.expand_dims([tokenizer.word_index['<start>']] * target.shape[0], 1)

    with tf.GradientTape() as tape:
        features = encoder(img_tensor)

        for i in range(1, target.shape[1]):
            # passing the features through the decoder
            predictions_val, hidden, _ = decoder(dec_input, features, hidden)
            loss_val += loss_func(target[:, i], predictions_val)
            accuracy_val += acc_func(target[:, i], predictions_val)

            # using teacher forcing
            dec_input = tf.expand_dims(target[:, i], 1)

    total_loss_val = (loss_val / int(target.shape[1]))
    total_accuracy_val = (accuracy_val / int(target.shape[1]))

```

```
return loss_val, total_loss_val, total_accuracy_val
```



```

In [ ]: EPOCHS = 50
val_loss_history = []
count = 0

for epoch in range(start_epoch, EPOCHS):
    print("=====Start Epoch" + str(epoch+1)+ "=====")

    total_loss_train = 0
    total_acc_train = 0
    total_loss_val = 0
    total_acc_val = 0
    print("Train loss")
    for (batch, (img_tensor, target)) in enumerate(dataset_train):
        batch_loss, t_loss, t_acc = train_step(img_tensor, target)
        total_loss_train += t_loss
        total_acc_train += t_acc

        if batch% 40 == 0:
            print('Epoch {} Batch{} Loss{:.4f} acc{:.4f}'.format(epoch+1, batch, batch_
with train_summary_writer.as_default():
            tf.summary.scalar('loss', total_loss_train/int(len(i_train)//BATCH_SIZE), step=
            tf.summary.scalar('accuracy', total_acc_train/int(len(i_train)//BATCH_SIZE), st

#Validation
print('Validation Loss')
for (batch, (img_tensor, target)) in enumerate(dataset_validation):
    batch_loss, v_loss, v_acc = val_step(img_tensor, target)
    total_loss_val += v_loss
    total_acc_val += v_acc
    if batch% 40 == 0:
        print('Epoch {} Batch {} Loss{:.4f} acc{:.4f}'.format(epoch+1, batch, batch_
with validation_summary_writer.as_default():
        tf.summary.scalar('loss', total_loss_val/int(len(i_validate)//BATCH_SIZE), step=
        tf.summary.scalar('accuracy', total_acc_val/int(len(i_validate)//BATCH_SIZE), s

print('EPOCH {}, LOSS: {}, ACCURACY:{}, TEST LOSS: {}, TEST ACCURACY: {}'.forma

val_loss_history.append(total_loss_val/int(len(i_validate)//BATCH_SIZE))
if epoch > 6:
    if count >= 4:
        print("Early Stopping Invoked!!!! Stopping the training \n\n")
        ckpt_manager.save()
        break
    else:
        if val_loss_history[epoch-1] < val_loss_history[epoch]:
            print("\n*****count++ Increased*****\n")
            count +=1

```

```

=====Start Epoch1=====
Train loss
Epoch 1 Batch0 Loss0.6172 acc0.0506

```

```

Epoch 1 Batch40 Loss0.4577 acc0.0520
Epoch 1 Batch80 Loss0.5037 acc0.0533
Validation Loss
Epoch 1 Batch 0 Loss0.4995 acc0.0536
EPOCH 1, LOSS: 0.5678637027740479, ACCURACY:5.274105072021484, TEST LOSS: 0.6
858891248703003, TEST ACCURACY: 5.631457805633545
=====Start Epoch2=====
Train loss
Epoch 2 Batch0 Loss0.4816 acc0.0542
Epoch 2 Batch40 Loss0.3206 acc0.0553
Epoch 2 Batch80 Loss0.5526 acc0.0562
Validation Loss
Epoch 2 Batch 0 Loss0.3706 acc0.0565
EPOCH 2, LOSS: 0.546222984790802, ACCURACY:5.5949859619140625, TEST LOSS: 0.6
736574172973633, TEST ACCURACY: 5.928240776062012
=====Start Epoch3=====

```

Evaluation

```

In [ ]: from tensorflow.keras.applications.densenet import preprocess_input
def get_img_tensor(image_path,img_name,model_image):
    img = tf.io.read_file(image_path+str(img_name))
    img = tf.image.decode_jpeg(img,channels=3)
    img = tf.image.resize(img,(224,224))
    img = preprocess_input(img)
    img_features = model_image(tf.constant(img)[None,:])
    return img_features

```

```

In [ ]: def evaluate(img_name):

    hidden = tf.zeros((1, units))
    img_tensor = tf.convert_to_tensor([get_img_tensor(img_path,img_name[0], image
                                                    get_img_tensor(img_path,img_name[1], image
    img_features = tf.constant(img_tensor)[None, :])
    features_val = encoder(img_features)

    dec_input = tf.expand_dims([tokenizer.word_index['<start>']], 0)
    result = []
    text = ""
    max_length_output=60
    for i in range(max_length_output):
        predictions, hidden,attention_weights = decoder(dec_input,features_val,hi
        predicted_id = tf.random.categorical(predictions,1)[0][0].numpy()
        if predicted_id ==0:
            word = ""
        else:
            word = tokenizer.index_word[predicted_id]
        result.append(word)
        text += " " + word
        if word == '<end>' or word == 'end':
            return result, text

        dec_input = tf.expand_dims([predicted_id], 0)
    return result, text


```

```

In [ ]: import matplotlib.image as mpimg
from nltk.translate.bleu_score import sentence_bleu
def test_img_cap(img_data, actual_text):
    result, text = evaluate(img_data)
    """Displays images for given input array of image names"""
    fig, axs = plt.subplots(1, len(img_data), figsize = (10,10), tight_layout=True)
    count = 0
    for img, subplot in zip(img_data, axs.flatten()):
        img=mpimg.imread(img_path+img)
        imgplot = axs[count].imshow(img_, cmap = 'bone')
        count +=1
    plt.show()
    reference = [actual_text.split()[1:-1]]
    result = result[:-1]
    print("="*50)
    print("Actual", actual_text)
    print("Predicted:",text)
    print("="*50)
    print('Individual 1-gram: {:.4f} Cumulative 1-gram: {:.4f}'.format(sentence_b
    print('Individual 2-gram: {:.4f} Cumulative 2-gram: {:.4f}'.format(sentence_b
    print('Individual 3-gram: {:.4f} Cumulative 3-gram: {:.4f}'.format(sentence_b
    print('Individual 4-gram: {:.4f} Cumulative 4-gram: {:.4f}'.format(sentence_b

```

```
In [ ]: test_img_cap(i_test[3],o_test[3])
```




```
=====
Actual <start> clear lungs no acute cardiopulmonary abnormality <end>
Predicted: no acute cardiopulmonary disease <end>
=====
```

```
Individual 1-gram: 0.4549 Cumulative 1-gram: 0.4549
Individual 2-gram: 0.4044 Cumulative 2-gram: 0.4289
Individual 3-gram: 0.3033 Cumulative 3-gram: 0.3839
Individual 4-gram: 0.6065 Cumulative 4-gram: 0.4289
```


```
/usr/local/lib/python3.6/dist-packages/nltk/translate/bleu_score.py:490: UserWarning:
Corpus/Sentence contains 0 counts of 4-gram overlaps.
BLEU scores might be undesirable; use SmoothingFunction().
  warnings.warn(_msg)
```

```
In [ ]: test_img_cap(i_test[64], o_test[64])
```



```
=====
Actual <start> no acute cardiopulmonary disease <end>
Predicted: mildly negative for todays study although this clips <end>
=====
Individual 1-gram: 0.0000 Cumulative 1-gram: 0.0000
Individual 2-gram: 0.0000 Cumulative 2-gram: 0.0000
Individual 3-gram: 0.0000 Cumulative 3-gram: 0.0000
Individual 4-gram: 0.0000 Cumulative 4-gram: 0.0000
```

```
In [ ]: test_img_cap(i_test[29], o_test[29])
```




```
=====
Actual <start> no acute cardiopulmonary disease <end>
Predicted:  no evidence of tuberculosis comparison chest <end>
=====
```

```
Individual 1-gram: 0.1667 Cumulative 1-gram: 0.1667
Individual 2-gram: 1.0000 Cumulative 2-gram: 0.4082
Individual 3-gram: 1.0000 Cumulative 3-gram: 0.5536
Individual 4-gram: 1.0000 Cumulative 4-gram: 0.6389
```

```
/usr/local/lib/python3.6/dist-packages/nltk/translate/bleu_score.py:490: UserWarning:
Corpus/Sentence contains 0 counts of 2-gram overlaps.
BLEU scores might be undesirable; use SmoothingFunction().
  warnings.warn(_msg)
```

```
In [ ]: test_img_cap(i_test[65], o_test[65])
```




```
=====
Actual <start> no acute intrathoracic abnormality <end>
Predicted: no suspicious chest prior available <end>
=====
```

```
Individual 1-gram: 0.2000 Cumulative 1-gram: 0.2000
Individual 2-gram: 1.0000 Cumulative 2-gram: 0.4472
Individual 3-gram: 1.0000 Cumulative 3-gram: 0.5879
Individual 4-gram: 1.0000 Cumulative 4-gram: 0.6687
```

```
/usr/local/lib/python3.6/dist-packages/nltk/translate/bleu_score.py:490: UserWarning:
Corpus/Sentence contains 0 counts of 2-gram overlaps.
BLEU scores might be undesirable; use SmoothingFunction().
  warnings.warn(_msg)
```

```
In [ ]: test_img_cap(i_test[103], o_test[103])
```




```
=====
Actual <start> no acute cardiopulmonary process <end>
Predicted: no acute cardiopulmonary abnormality <end>
=====
Individual 1-gram: 0.7500 Cumulative 1-gram: 0.7500
Individual 2-gram: 0.6667 Cumulative 2-gram: 0.7071
Individual 3-gram: 0.5000 Cumulative 3-gram: 0.6329
Individual 4-gram: 1.0000 Cumulative 4-gram: 0.7071
```

```
/usr/local/lib/python3.6/dist-packages/nltk/translate/bleu_score.py:490: User
Warning:
Corpus/Sentence contains 0 counts of 4-gram overlaps.
```



```
BLEU scores might be undesirable; use SmoothingFunction().
warnings.warn(_msg)
```

```
In [ ]: test_img_cap(i_test[363], o_test[363])
```




```
=====
Actual <start> no active disease <end>
Predicted: no acute pulmonary edema <end>
=====
```

```
Individual 1-gram: 0.2500 Cumulative 1-gram: 0.2500
Individual 2-gram: 1.0000 Cumulative 2-gram: 0.5000
Individual 3-gram: 1.0000 Cumulative 3-gram: 0.6329
Individual 4-gram: 1.0000 Cumulative 4-gram: 0.7071
```


```
/usr/local/lib/python3.6/dist-packages/nltk/translate/bleu_score.py:490: UserWarning:
Corpus/Sentence contains 0 counts of 2-gram overlaps.
BLEU scores might be undesirable; use SmoothingFunction().
warnings.warn(_msg)
```

```
In [ ]: test_img_cap(i_test[29], o_test[29])
```



```
=====
Actual <start> no acute cardiopulmonary disease <end>
Predicted: hypoinflation with the the lung nodules or scarring <end>
=====
Individual 1-gram: 0.0000 Cumulative 1-gram: 0.0000
Individual 2-gram: 0.0000 Cumulative 2-gram: 0.0000
Individual 3-gram: 0.0000 Cumulative 3-gram: 0.0000
Individual 4-gram: 0.0000 Cumulative 4-gram: 0.0000
```

```
In [ ]: test_img_cap(i_test[366], o_test[366])
```



```
=====
Actual <start> no acute cardiopulmonary abnormality <end>
Predicted:  no acute cardiopulmonary process <end>
=====
```

```
Individual 1-gram: 0.7500 Cumulative 1-gram: 0.7500
Individual 2-gram: 0.6667 Cumulative 2-gram: 0.7071
Individual 3-gram: 0.5000 Cumulative 3-gram: 0.6329
Individual 4-gram: 1.0000 Cumulative 4-gram: 0.7071
```

```
/usr/local/lib/python3.6/dist-packages/nltk/translate/bleu_score.py:490: UserWarning:
Corpus/Sentence contains 0 counts of 4-gram overlaps.
BLEU scores might be undesirable; use SmoothingFunction().
  warnings.warn(_msg)
```

```
In [ ]: columns = ["idx", "image_1", "image_2", "actual", "predicted", "score"]
df = pd.DataFrame(columns = columns)
for i in tqdm(range(len(i_test))):
    result, text_predicted = evaluate(i_test[i])
    actual = ' '.join([str(elem) for elem in o_test[i].split()[1:-1]])
    predicted = ' '.join([str(elem) for elem in result[1:]])
    df = df.append(pd.Series([i, i_test[i][0], i_test[i][1], actual, predicted, score]))
df.head(20)
```

```
1%|          | 3/399 [00:03<08:40, 1.31s/it]/usr/local/lib/python3.6/dist-packages/nltk/translate/bleu_score.py:490: UserWarning:
Corpus/Sentence contains 0 counts of 2-gram overlaps.
BLEU scores might be undesirable; use SmoothingFunction().
warnings.warn(_msg)
5%|█         | 18/399 [00:24<09:20, 1.47s/it]/usr/local/lib/python3.6/dist-packages/nltk/translate/bleu_score.py:490: UserWarning:
Corpus/Sentence contains 0 counts of 3-gram overlaps.
BLEU scores might be undesirable; use SmoothingFunction().
warnings.warn(_msg)
9%|██        | 35/399 [00:47<08:07, 1.34s/it]/usr/local/lib/python3.6/dist-packages/nltk/translate/bleu_score.py:490: UserWarning:
Corpus/Sentence contains 0 counts of 4-gram overlaps.
BLEU scores might be undesirable; use SmoothingFunction().
warnings.warn(_msg)
100%|██████████| 399/399 [08:54<00:00, 1.34s/it]
```

Out[169]:

	idx	image_1	image_2	actual	predicted	score
0	0	CXR1925_IM-0599-1001.png	CXR1925_IM-0599-3003.png	bihilar prominence may be secondary mild lymph...	active disease <end>	0
1	1	CXR2152_IM-0772-1001.png	CXR2152_IM-0772-2001.png	normal exam	active disease <end>	0
2	2	CXR1767_IM-0501-0001-0002.png	CXR1767_IM-0501-0001-0001.png	no acute cardiopulmonary abnormality	nodular opacities in the chest x the presence ...	0.266667
3	3	CXR3959_IM-2023-1001.png	CXR3959_IM-2023-3001.png	clear lungs no acute cardiopulmonary abnormality	limits no evidence of active disease <end>	0.142857
4	4	CXR742_IM-2298-2001.png	CXR742_IM-2298-1001.png	abnormal opacity in the right lung base which ...	comparison chest further examination of patien...	0.00293972
5	5	CXR2402_IM-0951-1001.png	CXR2402_IM-0951-2001.png	no acute cardiopulmonary findings	lung most notably in the left left mid lung ba...	0
6	6	CXR2318_IM-0891-2001.png	CXR2318_IM-0891-1001.png	no comparison chest x lungs lucency left chest...	acute pulmonary disease chronic liver postsurg...	0.022779
7	7	CXR1200_IM-0134-1001.png	CXR1200_IM-0134-2001.png	small right juxtahilar opacity may represent i...	acute cardiopulmonary abnormality <end>	0
8	8	CXR1388_IM-0246-1001.png	CXR1388_IM-0246-2001.png	low lung volumes otherwise clear	of the abdomen abdomen dr telephone <end>	0

	idx	image_1	image_2	actual	predicted	score
9	9	CXR2046_IM-0688-2002.png	CXR2046_IM-0688-1001.png	minimally displaced fracture involving the lat...	pulmonary disease <end>	0
10	10	CXR1935_IM-0605-0001-0002.png	CXR1935_IM-0605-0001-0001.png	right ij catheter tip in proximal right atrium...	wellexpanded and clear lungs <end>	0
11	11	CXR1964_IM-0629-1001.png	CXR1964_IM-0629-2001.png	no acute cardiopulmonary abnormality	focal atelectasis or a midthoracic vertebral d...	0.0243902
12	12	CXR190_IM-0583-1001.png	CXR190_IM-0583-3001.png	no acute disease	preoperative bowel pattern <end>	0
13	13	CXR2742_IM-1197-1001.png	CXR2742_IM-1197-2001.png	no acute findings	acute cardiopulmonary disease <end>	0.25
14	14	CXR1533_IM-0344-1001.png	CXR1533_IM-0344-1002.png	normal chest film	evidence for size and clear lungs <end>	0
15	15	CXR1263_IM-0179-3001.png	CXR1263_IM-0179-1001.png	no acute findings	acute cardiopulmonary abnormality <end>	0.25
16	16	CXR928_IM-2426-2002.png	CXR928_IM-2426-1001.png	no acute disease	acute cardiopulmonary process <end>	0.25
17	17	CXR3279_IM-1560-1001.png	CXR3279_IM-1560-1002.png	no acute cardiopulmonary abnormality	suspicious appearing lungs most could be nonco...	0
18	18	CXR536_IM-2143-1001.png	CXR536_IM-2143-2001.png	no acute radiographic cardiopulmonary process	evidence of normal mediastinal contour within ...	0.133333
19	19	CXR1534_IM-0345-4004.png	CXR1534_IM-0345-12001.png	no focal air space consolidation nodular opaci...	suspicious appearing interstitial markings wit...	0

```
In [ ]: from nltk.translate.bleu_score import corpus_bleu
actual_,predicted_ = list(),list()
for i in tqdm(range(len(i_test))):
    result, text_predicted = evaluate(i_test[i])
    actual = ' '.join([str(elem) for elem in o_test[i].split()[1:-1]])
    predicted = ' '.join([str(elem) for elem in result[1:]])
    actual_.append(actual)
    predicted_.append(predicted)
```

100%|██████████| 399/399 [01:50<00:00, 3.60it/s]

```
In [ ]: print('Corpus 1-gram: {:.4f}'.format(corpus_bleu(actual_,predicted_,weights=(1, 0, 0, 0)))
print('Corpus 2-gram: {:.4f}'.format(corpus_bleu(actual_, predicted_,weights=(0.5, 0.5, 0, 0)))
print('Corpus 3-gram: {:.4f}'.format(corpus_bleu(actual_,predicted_,weights=(0, 0, 1, 0)))
print('Corpus 4-gram: {:.4f}'.format(corpus_bleu(actual_,predicted_,weights=(0, 0, 0, 1)))
```

/usr/local/lib/python3.6/dist-packages/nltk/translate/bleu\_score.py:490: UserWarning:

Corpus/Sentence contains 0 counts of 2-gram overlaps.

BLEU scores might be undesirable; use SmoothingFunction().

warnings.warn(\_msg)

Corpus 1-gram: 0.1969

Corpus 2-gram: 0.4437

Corpus 3-gram: 1.0000

Corpus 4-gram: 1.0000

```
In [ ]: import pickle

with open('i_test.pkl','wb') as f:
    pickle.dump(i_test[103],f)
```

```
In [ ]: import pickle

with open('o_test.pkl','wb') as f:
    pickle.dump(o_test[103],f)
```

## Conclusion

- The model is built on a GRU model with Bahdanau Attention seems better model than the basic model.
- The model is able to generate different sentences which are similar to the original report but in some the model generates totally different sentences.
- We used densenet 121 to generate features from the images and pass it into encoder.
- We used a teacher forcing method for training the decoder model.
- In the end we used a greedy search for getting a new word index.

## Future Work

- We can modify the decoder layer with Bert Transformer inplace of the attention layer.
- We can increase the number of CNN layers to deep layer.

```
In [ ]:
```

