

```
In [ ]: cd /content/drive/My Drive/Xray report generation

/content/drive/My Drive/Xray report generation
```

```
In [ ]: ls

AttentionModel.ipynb          Medical.csv
Basic_Model.ipynb            model_1_plot.png
brucechou1983_CheXNet_Keras_0.3.0_weights.h5  NLMCXR_png/
Chexnetmodel/                'Old Files'/
'Copy of AttentionModel.ipynb'  preprocessed.csv
'Copy of SimpleModel.ipynb'    simple_encoder_decoder_plot.png
dup_preprocessed.csv          SimpleModel.ipynb
ecgen-radiology/              tokenizer1.pkl
image_feature_vector1.pkl     Untitled0.ipynb
indiana_projections.csv       Untitled1.ipynb
logs/
```

```
In [ ]: import re
import pandas as pd
import matplotlib.pyplot as plt
import string
from sklearn.utils import shuffle
from sklearn.model_selection import train_test_split
from tqdm import tqdm
import numpy as np
import tensorflow as tf
from nltk.translate.bleu_score import sentence_bleu
```

```
In [ ]: df = pd.read_csv('Medical.csv')
```

```
In [ ]: df.head()
```

```
Out[5]:
```

| | image_name | image_caption | comparison | indication | findings | impression |
|---|---|---------------------------|----------------------------------|--|---|--|
| 0 | CXR1_1_IM-0001-3001.png,CXR1_1_IM-0001-4001.png | xray chest pa and lateral | none | positive tb test | the cardiac silhouette and mediastinum size ar... | normal chest x |
| 1 | CXR10_IM-0002-2001.png,CXR10_IM-0002-1001.png | pa and lateral chest x | chest radiographs | male chest pain | the cardiomedial silhouette is within nor... | no acute cardiopulmonary process |
| 2 | CXR100_IM-0002-1001.png,CXR100_IM-0002-2001.png | chest v frontallateral pm | none | no indication | both lungs are clear and expanded heart and me... | no active disease |
| 3 | CXR1000_IM-0003-1001.png,CXR1000_IM-0003-2001.... | pa and lateral chest x | pa and lateral chest radiographs | male | there is increased opacity within the right up... | increased opacity in the right upper lobe with... |
| 4 | CXR1001_IM-0004-1001.png,CXR1001_IM-0004-1002.png | chest v frontallateral pm | none | dyspnea subjective fevers arthritis immigrant ... | interstitial markings are diffusely prominent ... | diffuse fibrosis no visible focal acute disease |

```
In [ ]: df.shape
```

```
Out[6]: (3851, 9)
```

```
In [ ]: df_projections = pd.read_csv('indiana_projections.csv')
```

```
In [ ]: df_projections.head()
```

```
Out[11]:
```

| | uid | filename | projection |
|---|-----|------------------------|------------|
| 0 | 1 | 1_IM-0001-4001.dcm.png | Frontal |
| 1 | 1 | 1_IM-0001-3001.dcm.png | Lateral |
| 2 | 2 | 2_IM-0652-1001.dcm.png | Frontal |
| 3 | 2 | 2_IM-0652-2001.dcm.png | Lateral |
| 4 | 3 | 3_IM-1384-1001.dcm.png | Frontal |

```
In [ ]: img_path = 'NLMCXR_png/'
```

```
In [ ]: # to find the matching pattern in a string - https://thepythonguru.com/python-regular-expressions/
def find_fr_lr_images(li):
    """
    This funct is used to find the lateral images and frontal images from our list
    """
    img_list = []
    last_imag = ""
    for img in li:
        projections = df_projections[df_projections['filename'].str.contains(re.search
        if 'Lateral'== projections:
            last_imag = img
        else:
            img_list.append(img)
    return img_list, last_imag
```

```
In [ ]: #Creating structured data from raw xml files
columns = ["image_1", "image_2", "impression"]
df_1 = pd.DataFrame(columns = columns)
columns = ["image_1", "image_2", "impression"]
df_dup = pd.DataFrame(columns = columns)
no_lateral = 0
for item in tqdm(df.iterrows()):
    l = item[1]['image_name'].split(',')
    if len(l) > 2:
        li, last_img = find_fr_lr_images(l)
        if last_img == "":
            no_lateral +=1
            li, last_img = li[:-1], li[-1]
        for i in li:
            image_1 = i
            image_2 = last_img
            df_1 = df_1.append(pd.Series([image_1, image_2,item[1]['impression']]))
    elif len(l) == 2:
        image_1 = l[0]
        image_2 = l[1]
        df_1 = df_1.append(pd.Series([image_1, image_2, item[1]['impression']], i
    elif len(l) == 1:
        #creating duplicate dataframe separately to keep it in all dataset train
        df_dup = df_dup.append(pd.Series([l[0], l[0],item[1]['impression']], inde
print("Total Report without Lateral images {}".format(no_lateral))
```

```
1155it [00:03, 312.59it/s]/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:10: DeprecationWarning: The truth value of an empty array is ambiguous. Returning False, but in future this will result in an error. Use `array.size > 0` to check that an array is not empty.
```

```
# Remove the CWD from sys.path while we load stuff.
3851it [00:13, 291.25it/s]
```

```
Total Report without Lateral images 1
```

```
In [ ]: df_1.shape
```

```
Out[15]: (3532, 3)
```

```
In [ ]: df_1.head()
```

```
Out[16]:
```

| | image_1 | image_2 | impression |
|---|--------------------------|--------------------------|---|
| 0 | CXR1_1_IM-0001-3001.png | CXR1_1_IM-0001-4001.png | normal chest x |
| 1 | CXR10_IM-0002-2001.png | CXR10_IM-0002-1001.png | no acute cardiopulmonary process |
| 2 | CXR100_IM-0002-1001.png | CXR100_IM-0002-2001.png | no active disease |
| 3 | CXR1000_IM-0003-1001.png | CXR1000_IM-0003-2001.png | increased opacity in the right upper lobe with... |
| 4 | CXR1000_IM-0003-3001.png | CXR1000_IM-0003-2001.png | increased opacity in the right upper lobe with... |

```
In [ ]: df_dup.shape
```

```
Out[17]: (446, 3)
```

```
In [ ]: df_dup.head()
```

```
Out[18]:
```

| | image_1 | image_2 | impression |
|---|--------------------------|--------------------------|---|
| 0 | CXR1003_IM-0005-2002.png | CXR1003_IM-0005-2002.png | retrocardiac soft tissue density the appearanc... |
| 1 | CXR1012_IM-0013-1001.png | CXR1012_IM-0013-1001.png | bibasilar airspace disease and bilateral pleur... |
| 2 | CXR1024_IM-0019-1001.png | CXR1024_IM-0019-1001.png | no acute abnormality |
| 3 | CXR1026_IM-0021-2002.png | CXR1026_IM-0021-2002.png | no acute cardiopulmonary disease |
| 4 | CXR1029_IM-0022-1001.png | CXR1029_IM-0022-1001.png | no pneumonia heart size normal scoliosis |

Adding Start and end tokens

```
In [ ]: def add_start_end_token(data):
    # Combining all the above students
    preprocessed_reviews_eng = []

    # tqdm is for printing the status bar
    for sentence in tqdm(data.values):
        sentence = '<start> ' + sentence + ' <end>'
        preprocessed_reviews_eng.append(sentence.strip())
    return preprocessed_reviews_eng
```

```
In [ ]: df_1['impression'] = add_start_end_token(df_1['impression'].astype(str))
```

100%|██████████| 3532/3532 [00:00<00:00, 771697.75it/s]

```
In [ ]: df_dup['impression'] = add_start_end_token(df_dup['impression'].astype(str))
```

100%|██████████| 446/446 [00:00<00:00, 396242.23it/s]

```
In [ ]: df_1.head()
```

```
Out[22]:
```

| | image_1 | image_2 | impression |
|---|--------------------------|--------------------------|---|
| 0 | CXR1_1_IM-0001-3001.png | CXR1_1_IM-0001-4001.png | <start> normal chest x <end> |
| 1 | CXR10_IM-0002-2001.png | CXR10_IM-0002-1001.png | <start> no acute cardiopulmonary process <end> |
| 2 | CXR100_IM-0002-1001.png | CXR100_IM-0002-2001.png | <start> no active disease <end> |
| 3 | CXR1000_IM-0003-1001.png | CXR1000_IM-0003-2001.png | <start> increased opacity in the right upper l... |
| 4 | CXR1000_IM-0003-3001.png | CXR1000_IM-0003-2001.png | <start> increased opacity in the right upper l... |

```
In [ ]: df_dup.head()
```

```
Out[23]:
```

| | image_1 | image_2 | impression |
|---|--------------------------|--------------------------|--|
| 0 | CXR1003_IM-0005-2002.png | CXR1003_IM-0005-2002.png | <start> retrocardiac soft tissue density the a... |
| 1 | CXR1012_IM-0013-1001.png | CXR1012_IM-0013-1001.png | <start> bibasilar airspace disease and bilater... |
| 2 | CXR1024_IM-0019-1001.png | CXR1024_IM-0019-1001.png | <start> no acute abnormality <end> |
| 3 | CXR1026_IM-0021-2002.png | CXR1026_IM-0021-2002.png | <start> no acute cardiopulmonary disease <end> |
| 4 | CXR1029_IM-0022-1001.png | CXR1029_IM-0022-1001.png | <start> no pneumonia heart size normal scolios... |

```
In [ ]: df_1.to_csv('preprocessed.csv')
```

```
In [ ]: df_dup.to_csv('dup_preprocessed.csv')
```

```
In [ ]: df_1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3532 entries, 0 to 3531
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   image_1     3532 non-null   object
1   image_2     3532 non-null   object
2   impression  3532 non-null   object
dtypes: object(3)
memory usage: 82.9+ KB
```

```
In [ ]: df_dup.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 446 entries, 0 to 445
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   image_1     446 non-null   object
1   image_2     446 non-null   object
2   impression  446 non-null   object
dtypes: object(3)
memory usage: 10.6+ KB
```

```
In [ ]: df_1 = pd.read_csv('preprocessed.csv')
```

```
In [ ]: df_1.shape
```

```
Out[9]: (3532, 4)
```

```
In [ ]: df_1.head()
```

```
Out[10]:
```

| | Unnamed: 0 | image_1 | image_2 | impression |
|---|------------|--------------------------|--------------------------|---|
| 0 | 0 | CXR1_1_IM-0001-3001.png | CXR1_1_IM-0001-4001.png | <start> normal chest x <end> |
| 1 | 1 | CXR10_IM-0002-2001.png | CXR10_IM-0002-1001.png | <start> no acute cardiopulmonary process <end> |
| 2 | 2 | CXR100_IM-0002-1001.png | CXR100_IM-0002-2001.png | <start> no active disease <end> |
| 3 | 3 | CXR1000_IM-0003-1001.png | CXR1000_IM-0003-2001.png | <start> increased opacity in the right upper l... |
| 4 | 4 | CXR1000_IM-0003-3001.png | CXR1000_IM-0003-2001.png | <start> increased opacity in the right upper l... |

```
In [ ]: df_2 = pd.read_csv("dup_preprocessed.csv")
```

```
In [ ]: df_2.head()
```

```
Out[12]:
```

| | Unnamed: 0 | image_1 | image_2 | impression |
|---|---------------|--------------------------|--------------------------|---|
| 0 | 0 | CXR1003_IM-0005-2002.png | CXR1003_IM-0005-2002.png | <start> retrocardiac soft tissue density the a... |
| 1 | 1 | CXR1012_IM-0013-1001.png | CXR1012_IM-0013-1001.png | <start> bibasilar airspace disease and bilater... |
| 2 | 2 | CXR1024_IM-0019-1001.png | CXR1024_IM-0019-1001.png | <start> no acute abnormality <end> |
| 3 | 3 | CXR1026_IM-0021-2002.png | CXR1026_IM-0021-2002.png | <start> no acute cardiopulmonary disease <end> |
| 4 | 4 | CXR1029_IM-0022-1001.png | CXR1029_IM-0022-1001.png | <start> no pneumonia heart size normal scolios... |

```
In [ ]: image_names = []
for img in tqdm(df['image_name'].str.split(',')):
    for i in range(len(img)):
        image_names.append(img[i])
```

100%|██████████| 3851/3851 [00:00<00:00, 894713.60it/s]

```
In [ ]: image_names[0:4]
```

```
Out[14]: ['CXR1_1_IM-0001-3001.png',
'CXR1_1_IM-0001-4001.png',
'CXR10_IM-0002-2001.png',
'CXR10_IM-0002-1001.png']
```

```
In [ ]: import tensorflow as tf
from tensorflow.keras.applications import densenet
from tensorflow.keras.applications.densenet import preprocess_input
from tensorflow.keras.layers import Dense, Dropout, Input, Conv2D
from tensorflow.keras.models import Model
```

```
In [ ]: chex = densenet.DenseNet121(include_top=False, weights=None, input_shape=(224, 224, 3)
X = chex.output
X = Dense(14, activation='sigmoid', name='predictions')(X)
model = Model(inputs=chex.input, outputs=X)
model.load_weights('brucechou1983_CheXNet_Keras_0.3.0_weights.h5')
image_features_model = Model(inputs=model.input, outputs=model.layers[-2].output)
```

```
In [ ]: # image_features_model.summary()
```

```
In [ ]: ## the input layer shape is (None,None,None,3)
        ## the output layer is (None,2048)
```

```
In [ ]: ##### This cell takes around 2 hrs to run.
        # image_feature_vectors = []
        # for img in tqdm(image_names):
        #     image = tf.io.read_file(img_path+str(img))
        #     image = tf.image.decode_jpeg(image,channels=3)
        #     image = tf.image.resize(image,(299,299))
        #     image = preprocess_input(image)
        #     image_features = image_features_model(tf.constant(image)[None,:]) # here we a
        #     image_features = tf.reshape(image_features,[-1,image_features.shape[1]])
        #     image_feature_vectors.append(image_features)
```

```
In [ ]: # print(image_feature_vectors[1])
```

```
In [ ]: # len(image_feature_vectors)
```

```
In [ ]: # import pickle

        # with open('image_feature_vector.pkl','wb') as f:
        #     pickle.dump(image_feature_vectors,f)
```

```
In [ ]: import pickle
        with open('image_feature_vector1.pkl','rb') as f:
            image_feature_vector_1 = pickle.load(f)
```

```
In [ ]: print(image_feature_vector_1[1])
```

```
tf.Tensor(
[[5.0062913e-04 2.0340595e-03 1.2797897e-03 ... 9.1456938e-01
  9.2038012e-01 7.4285048e-01]], shape=(1, 1024), dtype=float32)
```

Train, Validation and Test split

```
In [ ]: i_train,input_test,o_train,output_test = train_test_split(df_1[['image_1','image_2'],
        input_train,input_validate,output_train,output_validate = train_test_split(i_train,
```

```
In [ ]: input_train.shape,input_validate.shape,input_test.shape,output_train.shape,output_test.shape
```

```
Out[20]: ((2542, 2), (636, 2), (354, 2), (2542,), (636,), (354,))
```

```
In [ ]: i_train_dup,input_test_dup,o_train_dup,output_test_dup = train_test_split(df_2[['image_1','image_2'],
        input_train_dup,input_validate_dup,output_train_dup,output_validate_dup = train_t
```



```
In [ ]: input_train_dup.shape,input_validate_dup.shape,input_test_dup.shape,output_train_
```

```
Out[22]: ((320, 2), (81, 2), (45, 2), (320,), (81,), (45,))
```

Appending the duplicate data and the original data

```
In [ ]: i_train = np.append(input_train,input_train_dup,axis=0)
o_train = np.append(output_train,output_train_dup,axis=0)
i_validate = np.append(input_validate,input_validate_dup,axis=0)
o_validate = np.append(output_validate,output_validate_dup,axis=0)
i_test = np.append(input_test,input_test_dup,axis=0)
o_test = np.append(output_test,output_test_dup,axis=0)
```

```
In [ ]: i_train[0]
```

```
Out[24]: array(['CXR914_IM-2417-1001.png', 'CXR914_IM-2417-3001.png'], dtype=object)
```

```
In [ ]: #https://datascience.stackexchange.com/questions/24511/why-should-the-data-be-shu
# Here we will be shuffling the data
for i in range(5):
    i_train,o_train = shuffle(i_train,o_train,random_state=15)
    i_validate,o_validate = shuffle(i_validate,o_validate,random_state=15)
    i_test,o_test = shuffle(i_test,o_test,random_state=15)
```

TEXT TOKENIZATION

```
In [ ]: from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

max_len_output = 60

tokenizer = Tokenizer(oov_token="<unk>", filters='!"#$%&()*+.,-/:;=?@[\\]^_`{|}~ ')
tokenizer.fit_on_texts(o_train)
text_train = tokenizer.texts_to_sequences(o_train)
text_test = tokenizer.texts_to_sequences(o_test)
text_val = tokenizer.texts_to_sequences(o_validate)
dictionary = tokenizer.word_index

word2idx = {}
idx2word = {}
for k, v in dictionary.items():
    word2idx[k] = v
    idx2word[v] = k
```

```
In [ ]: vocab_size = len(word2idx)+1
vocab_size
```

```
Out[27]: 1339
```

```
In [ ]: print("==== Top 6 Word and its Index ====")
        list(dictionary.items())[:6]
```

```
==== Top 6 Word and its Index ====
```

```
Out[28]: [('<unk>', 1),
          ('<start>', 2),
          ('<end>', 3),
          ('no', 4),
          ('acute', 5),
          ('cardiopulmonary', 6)]
```

```
In [ ]: ## Padding the text
        text_train_output = pad_sequences(text_train,maxlen=60,dtype='int32',padding='post')
        text_validation_output = pad_sequences(text_val,maxlen=60,dtype='int32',padding='post')
        text_test_output = pad_sequences(text_test,maxlen=60,dtype='int32',padding='post')
```

```
In [ ]: text_train_output.shape
```

```
Out[30]: (2862, 60)
```

Creating Tensorflow Dataset

```
In [ ]: # converting the numpy image files into tensorformats

        def multi_image(img,imp):
            """
            This function will be taking in the imgaes we have and converting into tensor format
            """
            return tf.convert_to_tensor([image_feature_vector_1[image_names.index(img[0]).decode('utf-8')]])
```

```
In [ ]: dataset_train = tf.data.Dataset.from_tensor_slices((i_train,text_train_output))
```

```
In [ ]: dataset_validation = tf.data.Dataset.from_tensor_slices((i_validate,text_validation_output))
```

```
In [ ]: dataset_train = dataset_train.map(lambda item1,item2: tf.numpy_function(multi_image,[item1],tf.float32))
        dataset_validation = dataset_validation.map(lambda item1,item2: tf.numpy_function(multi_image,[item1],tf.float32))
```

```
In [ ]: BATCH_SIZE = 32
        BUFFER_SIZE = 500
        embedding_dimension = 256
        units = 512 # output
```

```
In [ ]: # Shuffle and Batch
dataset_train = dataset_train.shuffle(BUFFER_SIZE).batch(BATCH_SIZE)
dataset_train = dataset_train.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)

#Shuffle and Batch
dataset_validation = dataset_validation.shuffle(BUFFER_SIZE).batch(BATCH_SIZE)
dataset_validation = dataset_validation.prefetch(buffer_size=tf.data.experimental
```

```
In [ ]: # input_shape is (2,1,1024)
# vocab_size = 1225
```

```
In [ ]: vocab_size = 1339
input_layer = tf.keras.layers.Input(shape=(2,1,1024))
encoder_concat = tf.keras.layers.Concatenate()([input_layer[:,0], input_layer[:,1])
encoder_out = tf.keras.layers.Dense(embedding_dimension)(encoder_concat)
input_layer_text = tf.keras.layers.Input(shape=(1,))
x = tf.keras.layers.Embedding(vocab_size, embedding_dimension)(input_layer_text)
x = tf.keras.layers.Concatenate()([x, encoder_out])
x = tf.keras.layers.LSTM(units)(x)
x = tf.keras.layers.Dense(vocab_size)(x)
model = tf.keras.Model([input_layer, input_layer_text],x)
model.summary()
```

Model: "functional_5"

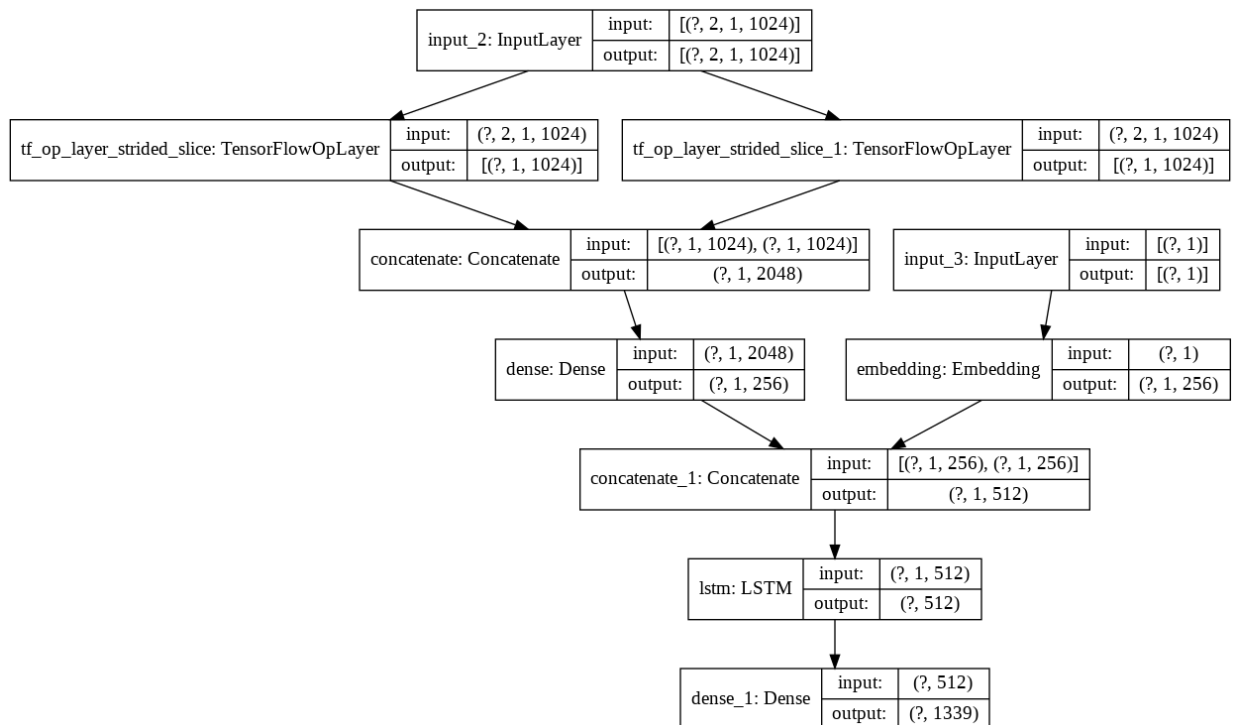
| Layer (type) | Output Shape | Param # | Connected to |
|---------------------------------|----------------------|---------|-------------------------------------|
| ===== | | | |
| input_2 (InputLayer) | [(None, 2, 1, 1024)] | 0 | |
| ===== | | | |
| tf_op_layer_strided_slice (Tens | [(None, 1, 1024)] | 0 | input_2[0][0] |
| ===== | | | |
| tf_op_layer_strided_slice_1 (Te | [(None, 1, 1024)] | 0 | input_2[0][0] |
| ===== | | | |
| input_3 (InputLayer) | [(None, 1)] | 0 | |
| ===== | | | |
| concatenate (Concatenate) | (None, 1, 2048) | 0 | tf_op_layer_st rided_slice[0][0] |
| ===== | | | |
| embedding (Embedding) | (None, 1, 256) | 342784 | input_3[0][0] |
| ===== | | | |
| dense (Dense) | (None, 1, 256) | 524544 | concatenate[0] [0] |
| ===== | | | |
| concatenate_1 (Concatenate) | (None, 1, 512) | 0 | embedding[0] dense[0][0] |
| ===== | | | |
| lstm (LSTM) | (None, 512) | 2099200 | concatenate_1 [0][0] |
| ===== | | | |
| dense_1 (Dense) | (None, 1339) | 686907 | lstm[0][0] |
| ===== | | | |
| Total params: 3,653,435 | | | |

Trainable params: 3,653,435

Non-trainable params: 0

```
In [ ]: tf.keras.utils.plot_model(model,to_file='simple_encoder_decoder_plot.png',show_sh
```

```
Out[40]:
```



```
In [ ]: class Encoder(tf.keras.Model):
    def __init__(self, embedding_dim):
        super(Encoder, self).__init__()
        self.fc = tf.keras.layers.Dense(embedding_dim, kernel_initializer=tf.kera
            name="encoder_output_layer")

    def call(self, x):
        x = tf.reshape(x, [x.shape[0], x.shape[1], x.shape[3]])
        encoder_concat = tf.keras.layers.Concatenate()([x[:,0], x[:,1]])
        x = self.fc(encoder_concat)
        x = tf.nn.relu(x)
        return x
```

```
In [ ]: class Decoder(tf.keras.Model):
    def __init__(self, embedding_dim, units, vocab_size):
        super(Decoder, self).__init__()
        self.units = units
        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
        self.lstm = tf.keras.layers.LSTM(self.units,
                                          return_sequences=True,
                                          return_state=True,
                                          recurrent_initializer=tf.keras.initializer
        self.dense = tf.keras.layers.Dense(vocab_size, kernel_initializer=tf.kera

    def call(self, x, features):
        #input x = input word teach forcing
        #input features = encoder image features
        x = self.embedding(x)
        x = tf.concat([x, tf.expand_dims(features,1)], axis=-1)
        output, state, _ = self.lstm(x)
        x = self.dense(output)
        return x
```

Encoder Decoder Model

```
In [ ]: # modelling loss and accuracy for gradient tape based on this https://www.tensorf

optimizer = tf.keras.optimizers.Adam()
loss_obj = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True, reduc
acc_obj = tf.keras.metrics.SparseCategoricalAccuracy()

def loss_func(real,pred):
    # getting the mask
    mask = tf.math.logical_not(tf.math.equal(real,0))

    # calculate the loss
    loss_ = loss_obj(real,pred)

    # converting mask dtype to loss dtype
    mask = tf.cast(mask,dtype=loss_.dtype)

    # applying the mask to loss
    loss_ = loss_ * mask
    loss_ = tf.reduce_mean(loss_)

    return loss_

def acc_func(real,pred):
    acc_f = acc_obj(real,pred)
    return tf.reduce_mean(acc_f)
```

```
In [ ]: import datetime
current_time = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
train_log_dir = 'logs/Basic_gradient_tape/' + current_time + '/train'
validation_log_dir = 'logs/Basic_gradient_tape/' + current_time + '/test'
train_summary_writer = tf.summary.create_file_writer(train_log_dir)
validation_summary_writer = tf.summary.create_file_writer(validation_log_dir)
```

```
In [ ]: encoder = Encoder(embedding_dimension)
decoder = Decoder(embedding_dimension, units, vocab_size)
```

```

In [ ]: @tf.function
def train_step(tensor, target):
    loss = 0
    accuracy = 0
    dec_input = tf.expand_dims([tokenizer.word_index['<start>']] * target.shape[0])
    with tf.GradientTape() as tape:
        features = encoder(tensor)
        for i in range(1, target.shape[1]):
            # passing the features through the decoder
            predictions = decoder(dec_input, features)
            loss += loss_func(target[:, i], predictions)
            accuracy += acc_func(target[:, i], predictions)

            # using teacher forcing
            dec_input = tf.expand_dims(target[:, i], 1)
            #print("decoder input teacher", dec_input.shape)
    total_loss = (loss / int(target.shape[1]))
    total_acc = (accuracy / int(target.shape[1]))
    trainable_variables = encoder.trainable_variables + decoder.trainable_variables

    gradients = tape.gradient(loss, trainable_variables)

    optimizer.apply_gradients(zip(gradients, trainable_variables))

    return loss, total_loss, total_acc

#validation function
@tf.function
def val_step(tensor, target):
    loss_val = 0
    accuracy_val = 0
    dec_input = tf.expand_dims([tokenizer.word_index['<start>']] * target.shape[0])
    with tf.GradientTape() as tape:
        features = encoder(tensor)
        for i in range(1, target.shape[1]):
            # passing the features through the decoder
            predictions_val = decoder(dec_input, features)
            loss_val += loss_func(target[:, i], predictions_val)
            accuracy_val += acc_func(target[:, i], predictions_val)

            # using teacher forcing
            dec_input = tf.expand_dims(target[:, i], 1)
            #print("decoder input teacher", dec_input)
    total_loss_val = (loss_val / int(target.shape[1]))
    total_acc_val = (accuracy_val / int(target.shape[1]))
    return loss_val, total_loss_val, total_acc_val

```

MODEL Training


```

In [ ]: tf.keras.backend.clear_session()
EPOCHS = 10
loss_plot_train = []
loss_plot_val = []
for epoch in range(0, EPOCHS):
    print("==== Start Epoch " +str(epoch + 1)+ " =====")

    total_loss_train = 0
    total_acc_train = 0
    total_loss_val = 0
    total_acc_val = 0
    print('Batchwise Train loss')
    for (batch, (jpg_tensor, target)) in enumerate(dataset_train):
        batch_loss, t_loss, t_acc = train_step(jpg_tensor, target)
        total_loss_train += t_loss
        total_acc_train += t_acc

        if batch % 40 == 0:
            print ('Epoch {} Batch {} Loss {:.4f} acc {:.4f}'.format(
                epoch + 1, batch, batch_loss / int(target.shape[1]), t_acc))

    loss_plot_train.append(total_loss_train / int(len(i_train) // BATCH_SIZE))
    with train_summary_writer.as_default():
        tf.summary.scalar('loss', total_loss_train/ int(len(i_train) // BATCH_SIZE))
        tf.summary.scalar('accuracy', total_acc_train/ int(len(i_train) // BATCH_SIZE))

    print('Batchwise validation loss')
    for (batch, (jpg_tensor, target)) in enumerate(dataset_validation):
        batch_loss_val, t_loss_val, t_acc_val = val_step(jpg_tensor, target)
        total_loss_val += t_loss_val
        total_acc_val += t_acc_val
        if batch % 40 == 0:
            print ('Epoch {} Batch {} Loss {:.4f} acc {:.4f}'.format(
                epoch + 1, batch, batch_loss_val / int(target.shape[1]), t_acc_val))

    with validation_summary_writer.as_default():
        tf.summary.scalar('loss', total_loss_val/int(len(i_validate) // BATCH_SIZE))
        tf.summary.scalar('accuracy', total_acc_val/int(len(i_validate) // BATCH_SIZE))

    template = 'Epoch {}, Loss: {}, Accuracy: {}, Test Loss: {}, Test Accuracy: {}'
    print (template.format(epoch+1,
        total_loss_train/ int(len(i_train) // BATCH_SIZE),
        (total_acc_train/ int(len(i_train) // BATCH_SIZE))*100,
        total_loss_val/int(len(i_validate) // BATCH_SIZE),
        (total_acc_val/int(len(i_validate) // BATCH_SIZE))*100))

```

==== Start Epoch 1 =====

Batchwise Train loss

Epoch 1 Batch 0 Loss 1.0213 acc 0.0004

Epoch 1 Batch 40 Loss 0.5824 acc 0.7893

Epoch 1 Batch 80 Loss 0.6593 acc 0.7930

Batchwise validation loss

Epoch 1 Batch 0 Loss 0.3677 acc 0.7922

Epoch 1, Loss: 0.6054882407188416, Accuracy: 76.15751647949219, Test Loss: 0.51

```
41300559043884, Test Accuracy: 82.40718078613281
===== Start Epoch 2 =====
Batchwise Train loss
Epoch 2 Batch 0 Loss 0.5559 acc 0.7840
Epoch 2 Batch 40 Loss 0.4391 acc 0.7782
Epoch 2 Batch 80 Loss 0.3942 acc 0.7741
Batchwise validation loss
Epoch 2 Batch 0 Loss 0.4417 acc 0.7728
Epoch 2, Loss: 0.4169798791408539, Accuracy: 78.60909271240234, Test Loss: 0.39
965489506721497, Test Accuracy: 80.63554382324219
===== Start Epoch 3 =====
Batchwise Train loss
Epoch 3 Batch 0 Loss 0.3044 acc 0.7703
Epoch 3 Batch 40 Loss 0.2925 acc 0.7681
Epoch 3 Batch 80 Loss 0.2865 acc 0.7662
Batchwise validation loss
Epoch 3 Batch 0 Loss 0.3120 acc 0.7661
Epoch 3, Loss: 0.33620578050613403, Accuracy: 77.6750259399414, Test Loss: 0.35
80491840839386, Test Accuracy: 80.03739929199219
===== Start Epoch 4 =====
Batchwise Train loss
Epoch 4 Batch 0 Loss 0.2925 acc 0.7647
Epoch 4 Batch 40 Loss 0.3061 acc 0.7643
Epoch 4 Batch 80 Loss 0.3507 acc 0.7629
Batchwise validation loss
Epoch 4 Batch 0 Loss 0.2451 acc 0.7629
Epoch 4, Loss: 0.3041955232620239, Accuracy: 77.23392486572266, Test Loss: 0.33
000192046165466, Test Accuracy: 79.68699645996094
===== Start Epoch 5 =====
Batchwise Train loss
Epoch 5 Batch 0 Loss 0.2505 acc 0.7619
Epoch 5 Batch 40 Loss 0.3086 acc 0.7613
Epoch 5 Batch 80 Loss 0.2852 acc 0.7609
Batchwise validation loss
Epoch 5 Batch 0 Loss 0.2808 acc 0.7608
Epoch 5, Loss: 0.2833741307258606, Accuracy: 76.99518585205078, Test Loss: 0.31
85734152793884, Test Accuracy: 79.48360443115234
===== Start Epoch 6 =====
Batchwise Train loss
Epoch 6 Batch 0 Loss 0.3131 acc 0.7600
Epoch 6 Batch 40 Loss 0.3046 acc 0.7601
Epoch 6 Batch 80 Loss 0.1760 acc 0.7595
Batchwise validation loss
Epoch 6 Batch 0 Loss 0.3143 acc 0.7594
Epoch 6, Loss: 0.27113163471221924, Accuracy: 76.8564453125, Test Loss: 0.31030
49397468567, Test Accuracy: 79.36990356445312
===== Start Epoch 7 =====
Batchwise Train loss
Epoch 7 Batch 0 Loss 0.3416 acc 0.7588
Epoch 7 Batch 40 Loss 0.3698 acc 0.7588
Epoch 7 Batch 80 Loss 0.2490 acc 0.7584
Batchwise validation loss
Epoch 7 Batch 0 Loss 0.3573 acc 0.7584
Epoch 7, Loss: 0.26025623083114624, Accuracy: 76.72484588623047, Test Loss: 0.3
112887442111969, Test Accuracy: 79.25816345214844
===== Start Epoch 8 =====
Batchwise Train loss
```

```

Epoch 8 Batch 0 Loss 0.1634 acc 0.7581
Epoch 8 Batch 40 Loss 0.2906 acc 0.7582
Epoch 8 Batch 80 Loss 0.3026 acc 0.7582
Batchwise validation loss
Epoch 8 Batch 0 Loss 0.4076 acc 0.7578
Epoch 8, Loss: 0.25657811760902405, Accuracy: 76.65787506103516, Test Loss: 0.3
164512813091278, Test Accuracy: 79.20201110839844
===== Start Epoch 9 =====
Batchwise Train loss
Epoch 9 Batch 0 Loss 0.2662 acc 0.7575
Epoch 9 Batch 40 Loss 0.2277 acc 0.7575
Epoch 9 Batch 80 Loss 0.1490 acc 0.7573
Batchwise validation loss
Epoch 9 Batch 0 Loss 0.3120 acc 0.7573
Epoch 9, Loss: 0.24627408385276794, Accuracy: 76.60527801513672, Test Loss: 0.3
0513450503349304, Test Accuracy: 79.16907501220703
===== Start Epoch 10 =====
Batchwise Train loss
Epoch 10 Batch 0 Loss 0.2724 acc 0.7570
Epoch 10 Batch 40 Loss 0.2121 acc 0.7570
Epoch 10 Batch 80 Loss 0.1736 acc 0.7567
Batchwise validation loss
Epoch 10 Batch 0 Loss 0.3214 acc 0.7568
Epoch 10, Loss: 0.24162065982818604, Accuracy: 76.54540252685547, Test Loss: 0.
2941944897174835, Test Accuracy: 79.09961700439453

```

```
In [ ]: %load_ext tensorboard
```

```
In [ ]: tensorboard --logdir=logs/
```

Output hidden; open in <https://colab.research.google.com> (<https://colab.research.google.com>) to view.

Evaluation

```

In [ ]: from tensorflow.keras.applications.densenet import preprocess_input
def get_img_tensor(image_path,img_name,model_image):
    img = tf.io.read_file(image_path+str(img_name))
    img = tf.image.decode_jpeg(img,channels=3)
    img = tf.image.resize(img,(224,224))
    img = preprocess_input(img)
    img_features = model_image(tf.constant(img)[None,:])
    return img_features

```

```

In [ ]: def evaluate(img_name):
    img_tensor = tf.convert_to_tensor([get_img_tensor(img_path,img_name[0], image_f
                                     get_img_tensor(img_path,img_name[1], image_f
    img_features = tf.constant(img_tensor)[None,:])
    feature_val = encoder(img_features)
    dec_input = tf.expand_dims([tokenizer.word_index['<start>']],1)
    result = []
    text = ""
    max_length_output = 60
    for i in range(max_length_output):
        predictions = decoder(dec_input,feature_val)# the output shape is 1X1X1225

        predictions = tf.reshape(predictions,[predictions.shape[0],predictions.shape[
        # print(predictions)
        # break

        predicted_id =tf.argmax(tf.math.log(predictions),1)[0].numpy()

        result.append(tokenizer.index_word[predicted_id])
        text += " " + tokenizer.index_word[predicted_id]
        if tokenizer.index_word[predicted_id] == '<end>':
            return result,text

        dec_input = tf.expand_dims([predicted_id],1)
    return result,text

```

```


In [ ]: import matplotlib.image as mpimg
from nltk.translate.bleu_score import sentence_bleu
def test_img_cap(img_data):
    result,text = evaluate(img_data)
    fig,axs = plt.subplots(1,len(img_data),figsize=(10,10),tight_layout=True)
    count = 0
    for img,subplot in zip(img_data,axs.flatten()):
        img_ = mpimg.imread(img_path+img)
        imgplot = axs[count].imshow(img_,cmap='bone')
        count +=1
    plt.show()

    print("Predicted:",text)

```

```
In [ ]: print("Actual",o_test[164])  
test_img_cap(i_test[164])
```

Actual <start> no acute pulmonary disease <end>




Predicted: no acute cardiopulmonary abnormality <end>

```
In [ ]:
```

```
In [ ]: print("Actual",o_test[64])  
        test_img_cap(i_test[64])
```

Actual <start> no acute cardiopulmonary disease <end>




Predicted: no acute cardiopulmonary abnormality <end>

```
In [ ]:
```

```
In [ ]: print("Actual",o_test[66])  
        test_img_cap(i_test[66])
```


Actual <start> heart size is normal and lungs are clear stable mm right midlung
perform granuloma <end>



Predicted: no acute cardiopulmonary abnormality <end>

```
In [ ]: print("Actual",o_test[266])  
test_img_cap(i_test[266])
```


Actual <start> no acute active cardiac pulmonary pleural disease <end>



Predicted: no acute cardiopulmonary abnormality <end>


```
In [ ]: print("Actual",o_test[29])  
        test_img_cap(i_test[29])
```

Actual <start> no acute cardiopulmonary disease <end>



Predicted: no acute cardiopulmonary abnormality <end>

```
In [ ]:
```